

# Notation from Yamasaki and Ushio explained

November 5, 2024

## Q Learning Lingo

- $\alpha$  is the *learning rate* or step size. According to Wikipedia, this variable determines to what extent newly acquired information overrides old information. A factor of 0 makes the agent learn nothing (exclusively exploiting prior knowledge). In contrast, a factor of 1 makes the agent consider only the most recent information (ignoring prior knowledge to explore possibilities).
- $\gamma$  is the *discounted rate of reward* or discount factor. Its value determines the importance of future rewards. A factor of 0 will make the agent “myopic” (or short-sighted) by only considering current rewards, while a factor approaching 1 will make it strive for a long-term high reward.

## Section 3 notation up to Eq. (13)

- $SV_i$  - the local supervisor, where  $i \in \{1, \dots, n\}$ 
  - I typically use the notation  $\mathcal{S}_i$  for a *supervisor*, which is another name for a *controller*, occasionally represented by  $\mathcal{C}_i$ . In truth, a while ago, I reverted to using  $H$  for the global (control) decision function and  $h_i$  for the local (control) decision functions, where  $i \in I$  and  $I$  is the index set of the local controllers (i.e.,  $I = \{1, \dots, n\}$ ).
- $G = (X, \Sigma, f, x_0)$  - the discrete-event system, where  $X$  is the finite set of states,  $\Sigma$  is the set of events,  $f$  is the state transition function and  $x_0$  is the initial state
  - I typically use  $G$  for the uncontrolled system, but I use  $Q$  as the state set, call  $\Sigma$  the alphabet, and use  $\delta : Q \times \Sigma \rightarrow Q$  as the (partial) transition function (recall that a complete function will have a transition of every element of  $\Sigma$  from every state in  $Q$ ); then  $q_0$  is the initial state.
- $\Sigma^c \subseteq \Sigma$  is the set of controllable events and  $\Sigma^o \subseteq \Sigma$  is the set of observable events
  - It is more typical to use  $\Sigma_c$  and  $\Sigma_o$  to denote the controllable and observable event sets, respectively.
  - Similarly, to denote the controllable and observable sets of events for a specific controller/supervisor/agent  $i \in I$ , you’ll see  $\Sigma_{c,i}$  and  $\Sigma_{o,i}$ .
- They use  $\epsilon$  to represent the empty string.
  - I prefer to use  $\varepsilon$ . It’s easier to distinguish from set membership  $\in$  depending on the font set used.
- $F_G(x)$  is used to denote the *active event set*. This is pretty antiquated vocabulary, although not unheard of depending on the research genealogy of the authors. This set simply denotes all the elements of  $\Sigma$  that are labels of outgoing transitions from  $x$ . As noted above, since  $\delta$  may not be a complete function, not all elements in  $\Sigma$  are involved in a transition from every state. So  $F_G(x)$  allows you to collect the events that “are defined” from state  $x$ . In modern terminology, we now call this the *feasible event set* and denote it by  $\Gamma$  in the centralized case, but by  $\Gamma_i$  for an observer  $\mathcal{O}_i$ , for  $i \in I$ .

- $M_i^e : \Sigma \rightarrow (\Sigma_i^o \cup \{\epsilon\})$  is the *natural projection* for supervisor/controller/agent  $i$ . In effect, this operator simply sets an event  $\sigma \in \Sigma$  to  $\epsilon$  if  $\sigma \notin \Sigma_i^o$ . The operator is extended to work on  $\Sigma^*$  and  $\Sigma_i^o$ . Think of it as an eraser that removes all events that an agent cannot observe.
  - I almost exclusively use  $\pi$  to represent natural projection and, thus,  $\pi_i$  for the projection for supervisor/controller/agent  $i \in I$
- $SV_i = (S_i, \Sigma_i^o, g_i, x_0)$  is used to denote a local controller (supervisor, agent). [I think there are typos here as it seems silly to use  $S_i$  to represent the state set and then  $x_0$  as the initial state.] Usually, this is what we’re trying to produce as output. But for the reinforcement learning application, it seems as if they start with each controller’s local view of  $G$  and then use the rewards/penalties to converge to a better solution.
  - The more standard notation is to invoke the observer  $\mathcal{O}$ , which follows the construction of the subset construction algorithm. I typically use  $\mathcal{O}_i = (X_i, \Sigma_{o,i}, \delta_i, x_{0,i})$ , where  $X_i \subseteq 2^Q$  (the power set of states in  $Q$ ).
- The paper talks about a *control pattern*, which simply refers to a function that defines, for each controller, the events that are enabled at each state. Note that by definition, any event in  $\Sigma \setminus \Sigma_c$  is called *uncontrollable* and is permanently enabled (i.e., cannot be prevented from happening). In this paper, the local control patterns are combined via intersection to determine the overall (aka *global*) control pattern at that state.

## Notation from Eq. (14) and onwards in Section 3

- $\mathcal{P}_i^1(s_i, \pi_i, \sigma)$  is the probability that supervisor/agent  $i$  *observes*  $\sigma \in \Sigma_{o,i} \cap \pi_i$  when it chooses the control pattern  $\pi_i$  at state  $s_i$ . It is computed in Eq. (16)
- $\mathcal{P}_i^2(s_i, \sigma, s'_i)$  is the probability that supervisor/agent  $i$  *takes* a transition from state  $s_i$  to state  $s'_i$  via event  $\sigma \in \Sigma_{o,i}$ .
- $\mathcal{P}_i(s_i, \pi_i, s'_i)$  is then the combination of these two probabilities, defined by Eq. (15), summed over agent  $i$ ’s observable events in  $\pi_i \cap \Sigma_{o,i}$

## Section 4 Algorithm notation

- $T_i(s_i, \sigma)$  is a learning parameter, which seems to be tied to  $T_i^*$  in Eq. (20), where it is referred to as a discounted expected total reward. The parameter itself is initialized according to Eq. (21). The authors note that this is the reward when agent  $i$  observes an event  $\sigma \in \Sigma_{o,i}$  at state  $s_i$
- $R_i^1$  is another learning parameter, which is introduced as part of Eq. (18). It is defined as the *expectation of a reward* when supervisor  $i$  selects control pattern  $\pi_i$ . The authors say that this corresponds to the cost of disabling controllable events that are not included in the control pattern (since it represents events that are enabled at a given state). This learning parameter uses  $\beta$ .
- $\eta_i$  is the final learning parameter. After reading the authors’ previous paper on a centralized controller and reinforcement learning, it seems that  $\eta_i(s, \sigma)$  is simply a number between 0 and 1 that is assigned to each observable event  $\sigma \in \Sigma_{o,i}$  such that for all feasible and observable events at state  $s_i \in S_i$ ,

$$\sum_{\sigma \in F(s) \cap \Sigma_{o,i}} \eta_i(s_i, \sigma) = 1.$$

This is “computed” on the automaton that defines  $SV_i$ . For example, in  $SV_1$  in state 14, there are two feasible and observable events:  $m2$  and  $c2$ . So one possible assignment could be  $\eta_1(14, m2) = 0.5$  and

$\eta_1(14, c2) = 0.5$  because their sum is 1.0. Similarly, you could (randomly) define  $\eta_1(14, m2) = 0.65$  and  $\eta_1(14, c2) = 0.35$ . And none of these individual values can be 0, since they are part of the feasible alphabet at state 14. The values assigned to the  $\eta_i$  parameter are used as a starting point to determine the probabilities of an event occurring, which is defined by Eq. (16) over the control patterns, not the feasible events. This learning parameter uses  $\delta$ .

## Values used in the Cat and Mouse Example

- All  $Q$  values were initialized to 0.
- $\alpha = 0.1$  (According to Wikipedia, this is a standard value)
- $\beta = 0.1$
- $\delta = 0.1$
- $\gamma = 0.9$

## The original algorithm

---

**Algorithm 1** Original decentralized reinforcement learning algorithm

---

```

1: Initialize  $T_i$ ,  $R_i^1$ ,  $\eta_i$ , and  $Q_i$  for all  $SV_i$ , where  $i \in \{1, \dots, n\}$ 
2:  $\triangleright$  The true system state  $s$  is observed a state estimate  $s_i$  for each of the  $i \in \{1, \dots, n\}$   $\triangleleft$ 
3: for each episode do
4:    $s \leftarrow x_0$ 
5:   for each  $SV_i$  do
6:      $\triangleright$  We want to initialize each supervisor's initial state to be the state estimate of  $x_0$ , the initial state of  $G$ . We have no notation for this right now.  $\triangleleft$ 
7:     Initialize a state  $s_i \leftarrow x_0$ 
8:      $atMarkedState \leftarrow \text{False}$ 
9:     Select a control pattern  $\pi_i(s_i) \in \Pi_i(s_i)$  based on the  $Q_i$  values by  $SV_i$ .
10:    Apply global control pattern  $\pi(s) \leftarrow \cap_{i \in I} \pi_i(s_i)$  to the system (aka DES)  $G$ 
11:    Choose an event  $\sigma \in \Gamma(s) \cap \pi(s)$  to occur in  $G$  (Guided by the use of Eq. (9))
12:     $\triangleright$  Update the current state of  $G$   $\triangleleft$ 
13:     $s \leftarrow \delta(s, \sigma)$ 
14:    for all  $SV_i$  such that  $\sigma \in \Sigma_{o,i}$  and not  $atMarkedState$  do
15:      Observe the occurrence of event  $\sigma \in \Sigma_{o,i}$ 
16:      Acquire rewards  $r_i^1$  and  $r_i^2$ .
17:      In  $SV_i$ , make a transition  $g_i(s_i, \sigma) = s'_i$ .
18:      Update  $T_i(s_i, \sigma)$ ,  $R_i^1(s_i, \pi_i)$  and  $\eta_i(s_i, \sigma)$  using Eqs. X, Y and Z, respectively.
19:      Update  $Q_i$  values by Eq. XX.
20:       $s_i \leftarrow s'_i$ 
21:      if  $s_i$  is a marked state then
22:         $atMarkedState \leftarrow \text{True}$ 

```

---

---

**Algorithm 2** Decentralized Reinforcement Learning algorithm

---

```

1: procedure A METHOD( $G, \mathcal{O}_1, \dots, \mathcal{O}_n$ )  $\triangleright \mathcal{O}_i = (X_i, \Sigma_{o,i}, \delta_i, x_{i,0})$ 
   $\triangleright$  The set of control patterns (typically called a control action) at a state  $x_i \in X_i$  is the product of (i)
   $\Sigma_{uc,i}$  intersected with the feasible events at  $x_i$ , that is  $\Gamma_i(x_i)$  and (ii) the elements of the power set of
   $\Sigma_{c,i}$  intersected with  $\Gamma_i(x_i)$ .
   $\triangleright h_i(x_i)$  is the local control action (aka the control pattern) for  $\mathcal{O}_i$  at state  $x_i$ . What is it initialized to?
  All events in  $\Gamma_i(x_i) \cup (\Sigma_{uc,i} \cap \Gamma_i(x_i))$ ?
2:   for  $i \in I$  do
3:     for  $x_i \in X_i$  do
4:        $\triangleright T_i(x_i, \sigma_i), \eta_i(x_i, \sigma_i)$   $\triangleleft$ 
5:        $T_i \leftarrow 0$ 
6:        $\eta_i \leftarrow$ 
7:        $\triangleright R_i^1(x_i, h_i(x_i))$   $\triangleleft$ 
8:        $R_i^1 \leftarrow 0$ 
9:        $Q_i(x_i, h_i(x_i)) \leftarrow 0$ 
10:     $\triangleright$  End of initialization (Lines 1 and 2 of the original paper)  $\triangleleft$ 
11:     $\triangleright$  Lines below form the basis of an “episode” or epoch  $\triangleleft$ 
12:     $atMarkedState \leftarrow \text{False}$ 
13:    while not  $atMarkedState$  do
14:       $q \leftarrow q_0$   $\triangleright$  initialize current state of  $Q$ 
15:      for  $i \in I$  do
16:         $x_i \leftarrow x_{i,0}$   $\triangleright$  initialize current state of each  $\mathcal{O}_i$ 
17:        Choose the “best”  $h_i(x_i)$  according to current  $Q_i$  values for  $\mathcal{O}_i$ 
18:        Apply  $H(q) \leftarrow \cap_{i \in I} h_i(x_i)$  to  $G$ 
19:        Choose  $\sigma \in (\Gamma(q) \cap H(q))$  guided by Eq. (9)
20:         $q \leftarrow \delta(q, \sigma)$   $\triangleright$  Update current state of  $G$ 
21:        for  $i \in I$  such that  $\sigma \in \Sigma_{o,i}$  do
22:          Observe  $\sigma_i \in \Gamma_i(x_i)$ 
23:          Acquire rewards  $r_i^1$  and  $r_i^2$ 
24:          Update  $T_i(x_i, \sigma_i)$ 
25:          Update  $\eta_i(x_i, \sigma_i)$ 
26:          Update  $R_i^1(x_i, h_i(x_i))$ 
27:          Update  $Q_i$  according to Equation 24
28:           $x_i \leftarrow \delta_i(x_i, \sigma_i)$ 
29:          if  $x_i \in Marked(X_i)$  then
30:             $atMarkedState \leftarrow \text{True}$ 

```

---